

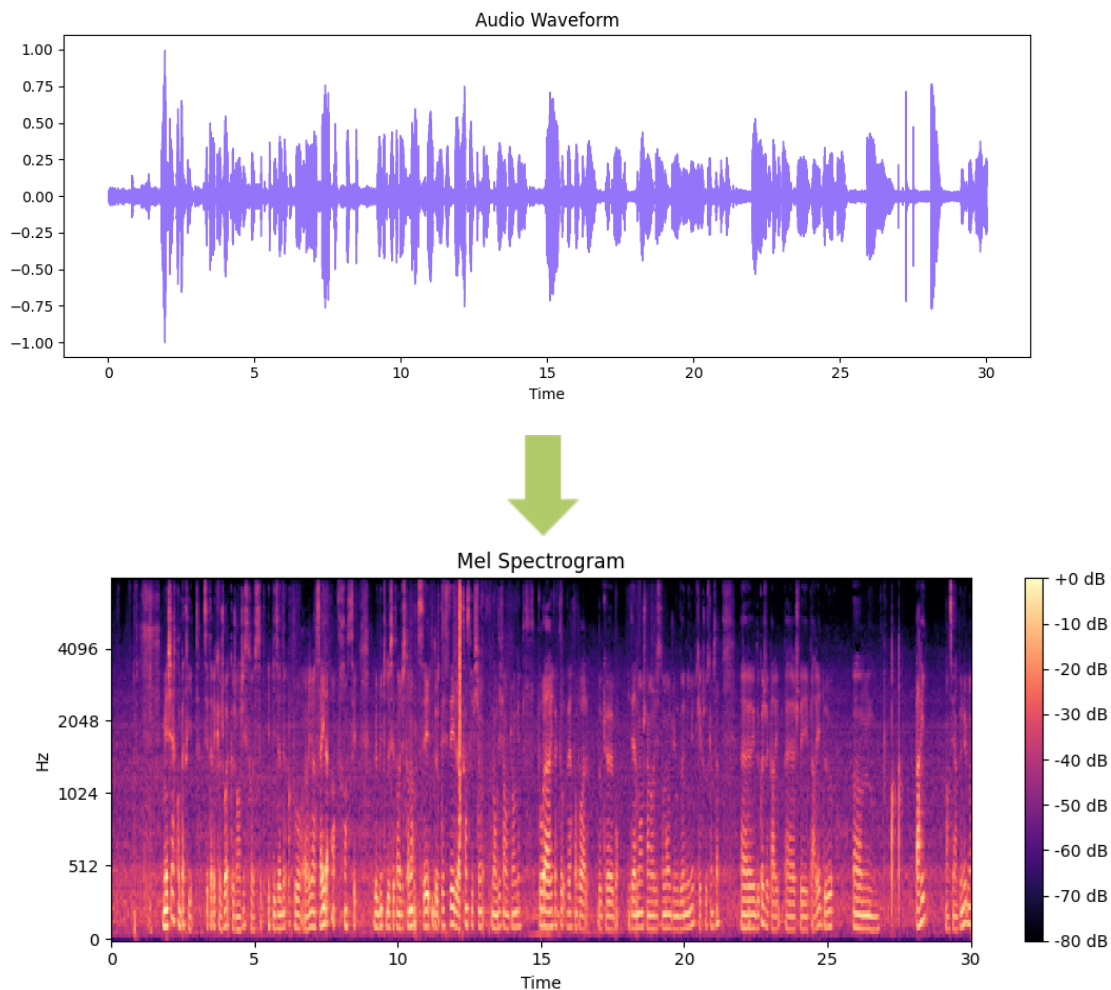
# Riverside.fm - Home assignment

## Product context

At Riverside, we use a speech-to-text model to extract the transcription from a given recording. The speech-to-text model requires pre-processing the recording into Mel-spectrograms.

*What is a Mel-spectrogram ?*

A Mel-spectrogram is a spectral representation of an audio signal on a Mel scale. You don't need to know what it is to complete this assignment, but this is intended to give more insight on the data we use.



The speech-to-text model expects a Mel-spectrogram of shape  $80 \times 3000$  (80 bins, 100 points per second). If the audio is less than 30 seconds, the spectrogram is padded with zeros to ensure shape consistency.

You can read more about mel-spectrograms [here](#).

These inputs are put in a queue and we want to batch them for optimization purposes.

# BatchQueue

In this assignment you will implement a new class for pytorch called: **BatchQueue**.

This queue differs from a standard queue in the fact that it supports a 'batch' of different queues with a **single tensor**. If the dimensions of a standard queue are  $[L, F]$  with  $L$  representing the length of the queue and  $F$  the number of features, then this queue will have an extra dimension and will have dims  $[L, B, F]$  with  $B$  representing the batch size.

In our case,  $F$  is the number of spectrogram features which we have flatten for ease of use, which means  $F=240000$  ( $80 \times 3000$ ).

We have provided a file `mel_spectrograms.npy` that has a shape of  $(5, 4, 240000)$  – it is a numpy array of 20 flattened spectrograms, reshaped as 5 batches of 4 elements. This file is used in the provided `sanity_check()` method to verify your code.

*Please implement the BatchQueue class either in the `batch_queue.py` file or the Colab notebook `BatchQueue.ipynb`. The Colab notebook already has pytorch installed.*

**Optimized performance is very important, make sure your implementations is efficient.**

The queue needs to support the following methods:

1. `init(dims):`
  - a. Arguments:
    - i. `dims` : Tensor = 1D tensor of size (3) specifying the dimensions of the queue:  $[L, B, F]$ . The dimensions should be treated as follows:
      1.  $L$  - queue length
      2.  $B$  - batch size
      3.  $F$  - number of features
  - b. Does not return anything
2. `dequeue(batch_indices):`
  - a. Arguments:
    - i. `batch_indices` : Tensor = 1D tensor of Integers of size  $(N)$ ,  $1 \leq N \leq B$
  - b. Returns a 2D tensor of size  $(N, F)$  representing the values in the head of the queues of indices `batch_indices`. Modifies the queues in indices `batch_indices` such that each instance is 'moved forward' in line. You may assume the queues are not empty.
3. `enqueue(values, batch_indices):`
  - a. Arguments:
    - i. `Values` : Tensor = 3D tensor of size  $(T, N, F)$ ,  $1 \leq T \leq L$ ,  $1 \leq N \leq B$  to be inserted at the tail of the queue of indices `[batch_indices]`. You may assume the queue has room for the values given.
    - ii. `Batch_indices` : Tensor = 1D tensor of Integers of size  $N$ ,  $1 \leq N \leq B$
  - b. Does not return anything
4. `peek(location, batch_indices):`
  - c. Arguments:
    - i. `Location` : Tensor = list of Strings of length  $(N)$ ,  $1 \leq N \leq B$  that can either be 'head' or 'tail'.
    - ii. `Batch_indices` : Tensor = 1D tensor of Integers of size  $(N)$ ,  $1 \leq N \leq B$
  - d. Returns a 2D tensor of size  $(N, F)$  representing the values in the head/tail of the queues at indices `batch_indices`. You may assume the queues are not empty.

# Visual Example

**dequeue(batch\_indices):**

- Batch\_queue size=(L=4,B=3,F)
- N=2, batch\_indices=torch.tensor([1,2])

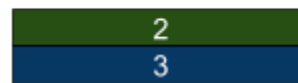


batch\_queue  
size=(L=4,B=3,F)

dequeue(batch\_indices)



batch\_queue  
size=(L=4,B=3,F)



Return tensor  
size=(N=2,F)